

Question 3c

- Product(maker, model, price)
- PC(model, speed)
- Printer(model, type)

2

Write in SQL: For each maker, find the minimum and maximum price of a (PC, ink-jet printer) combination.

```
SELECT p1.maker, min(p1.price+p2.price), max(p1.price+p2.price)
FROM Product p1, Product p2, PC pc, Printer t
WHERE t.type = 'ink-jet' and p1.model = pc.model and p2.model =
      t.model and p1.maker=p2.maker
GROUP BY p1.maker
```

2

Question 4

3

Given the instance of two relations:

| R: | A | B |
|----|---|---|
| | 1 | 2 |
| | 3 | 4 |
| | 1 | 3 |

| S: | B | C |
|----|---|---|
| | 1 | 3 |
| | 2 | 4 |

a) What is the result of the following query:

```
SELECT DISTINCT R.A
FROM R
WHERE R.A NOT IN (SELECT DISTINCT S.B AS A
                  FROM S
                  WHERE S.B = S.C)
```

A

1

3

3

Question 4b

4

| R: | A | B |
|----|---|---|
| | 1 | 2 |
| | 3 | 4 |
| | 1 | 3 |

| S: | B | C |
|----|---|---|
| | 1 | 3 |
| | 2 | 4 |

```
SELECT R.A, S.Cavg(R.B) as av
FROM R, S
WHERE R.B < 4
GROUP BY R.A, S.C
HAVING max(R.B) >= 2
```

| A | R.B | S.B | C |
|---|-----|-----|---|
| 1 | 2 | 1 | 3 |
| 1 | 2 | 2 | 4 |
| 3 | 4 | 1 | 3 |
| 3 | 4 | 2 | 4 |
| 1 | 3 | 1 | 3 |
| 1 | 3 | 2 | 4 |

| A | R.B | S.B | C |
|---|-----|-----|---|
| 1 | 2 | 1 | 3 |
| 1 | 3 | 1 | 3 |
| 1 | 2 | 2 | 4 |
| 1 | 3 | 2 | 4 |

| A | C | av |
|---|---|-----|
| 1 | 3 | 2.5 |
| 1 | 4 | 2.5 |

4

Question 5

5

Consider the following CREATE TABLE definition:

```
CREATE TABLE Midterm
(A INT NOT NULL,
 B INT NOT NULL,
 C INT NOT NULL,
 PRIMARY KEY (A),
 FOREIGN KEY (B) REFERENCES Midterm(A) ON DELETE CASCADE ON UPDATE CASCADE,
 FOREIGN KEY (C) REFERENCES Midterm(A) ON DELETE CASCADE ON UPDATE RESTRICT)
```

Consider the following instance table Midterm:

| A | B | C |
|---|---|---|
| 4 | 3 | 3 |
| 3 | 4 | 3 |

a) What is the result of the following statement:

```
UPDATE Midterm
SET B = B+1
WHERE B in (SELECT A FROM Midterm)
```

Answer: Error, the FK constraint is violated

5

6

End

6

B+ Tree Index

- The B+ tree structure is the most common index type in databases
- Index files can be quite large, often stored on disk, partially loaded into memory as needed
- Each node is at least 50% full

Level 1: [30, 65] (root)

Level 2: [15, 21], [41], [78] (Internal (index) nodes)

Level 3: [2, 7, 16, 20], [22, 29, 35, 37, 42, 50], [67, 77, 80, 91] (leaf nodes (data entries))

Credit: S. Lee

7

B+ Tree Index

Supports equality and range-searches efficiently

Non-leaf Pages (direct search)

Leaf Pages (Sorted by search key)

index entry

| | | | | | | | |
|----------------|----------------|----------------|----------------|----------------|-----|----------------|----------------|
| P ₀ | K ₁ | P ₁ | K ₂ | P ₂ | ... | K _m | P _m |
|----------------|----------------|----------------|----------------|----------------|-----|----------------|----------------|

Credit: Renee Miller

8

Example

Root: [17]

Internal nodes: [5, 13], [27, 30]

Leaf nodes: [2*, 3*, 5*, 7*, 8*, 14*, 16*, 22*, 24*, 27*, 29*, 33*, 34*, 38*, 39*]

- Find 28*? 29*? All > 15* and < 30*
- Insert/delete: Find data entry in leaf, then change it. Need to adjust parent sometimes.
- And change sometimes bubbles up the tree


9

Inserting a Data Entry

10

- ❑ Find correct leaf L.
- ❑ Put data entry onto L.
 - ❑ If L has enough space, done!
 - ❑ Else, must **split** L (into L and a new node L2)
 - ❑ Redistribute entries evenly, **copy up** middle key.
 - ❑ Insert index entry pointing to L2 into parent of L.
- ❑ This can happen recursively
 - ❑ To split index node, redistribute entries evenly, but **push up** middle key.
- ❑ Splits “grow” tree; root split increases height.

Insert data
value 4




10

Deleting a Data Entry

11

- ❑ Start at root, find leaf L where entry belongs.
- ❑ Remove the entry.
 - ❑ If L is at least half-full, done!
 - ❑ If not,
 - ❑ Try to **re-distribute**, borrowing from **sibling** (adjacent node with same parent as L).
 - ❑ If re-distribution fails, **merge** L and sibling.
- ❑ If merge occurred, must delete entry (pointing to L or sibling) from parent of L.
- ❑ Merge could propagate to root, decreasing height.

Delete value 3




11

Balanced vs. Unbalanced Trees


12

- ❑ In a balanced tree, every path from the root to a leaf node is the same length.

o Balanced



o Unbalanced



Credit: S. Lee

12

Hash Based Indexes

13

- ❑ Good for equality selections.
- ❑ Index is a collection of **buckets**
 - ❑ Bucket = **primary page** plus zero or more **overflow pages**.
 - ❑ Buckets contain data entries.
- ❑ **Hashing function h**: $h(r)$ = bucket in which (data entry for) record r belongs. h looks at the **search key** fields of r.
 - ❑ No need for “index entries” in this scheme.

13