

Assignment 3

16

- Posted on Avenue and MS Teams
- Due: December 3, 2021
- Recommend you start early, don't leave it until the last minute.

16

Scheduling Transactions

17

- Serial schedule: Schedule that does not interleave the actions of different transactions.
- Equivalent schedules: For any database state, the effect (on the set of objects in the database) of executing the first schedule is identical to the effect of executing the second schedule.
- Serializable schedule: A schedule that is equivalent to some serial execution of the transactions.

17

Serial Schedule

18

T1	T2
R(A)	
A := A+100	
R(B)	
B := B+100	
	R(A)
	A := A*2
	R(B)
	B := B*2

S: R1(A), W1(A), R1(B), W1(B), R2(A), W2(A), R2(B), W2(B)

T1
T2

Adapted from M. Balazinska

18

A Serializable Schedule

19

T1	T2
R(A)	
A := A+100	
	R(A)
	A := A*2
R(B)	
B := B+100	
	R(B)
	B := B*2

Notice: this is not a serial schedule, i.e., there is interleaving of operations

Net effect is the same as the serial schedule

S: R1(A), W1(A), R2(A), W2(A), R1(B), W1(B), R2(B), W2(B)

19

A Non-Serializable Schedule

20

T1	T2
R(A)	
A := A+100	
	R(A)
	A := A* 2
	R(B)
	B := B*2
R(B)	
B := B+100	

What's different?

Ordering of operations is not consistent

S: R1(A),W1(A), R2(A), W2(A), R2(B), W2(B), R1(B), W1(B)

20

Conflict operations

21

- Two operations in a schedule are said to be in *conflict* if they satisfy all three of the following conditions:

- (1) They belong to different transactions
- (2) They access the same item A;
- (3) At least one of the operations is a write(A)

Example in Sa: R1(A), R2(A), W1(A), W2(A), A1, C2;

- R1(A),W2(A) conflict, so do R2(A),W1(A),
- R1(A), W1(A) do not conflict because they belong to the same transaction,
- R1(A),R2(A) do not conflict because they are both read operations.

21

Conflicts

22

What can go wrong:


- Reading uncommitted data (WR), aka "dirty reads"
- Unrepeatable reads (RW)
- Lost updates (WW)

22

Reading Uncommitted Data

23

- WR Conflicts, "dirty reads"



- What's the problem?


T2 has read an updated value A, which is later aborted!

23

Unrepeatable Reads

24

RW Conflicts



What's the problem?

T2 has changed A that has been read by T1, while T1 is still in progress


24

Lost Updates

25

Overwriting uncommitted data

WW Conflicts



What's the problem?

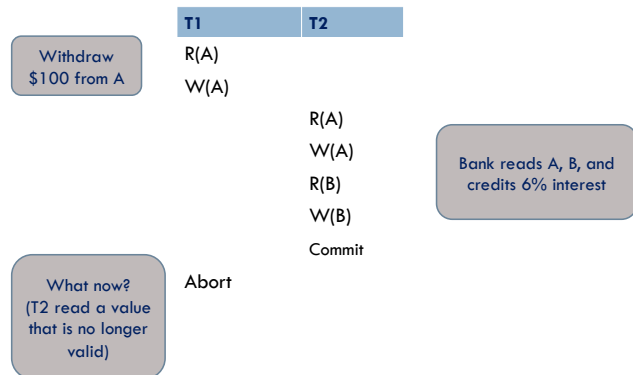
T2 overwrites T1's update of A, while T1 is still in progress.

T1: sets salaries to \$1000
T2: sets salaries to \$2000
Require: A and B salaries to be equal

25

Schedules with Aborts

26



26

Options

27

- If T2 did not commit, we abort T1 and **cascade** to T2.

- But T2 committed, so we cannot undo

- This schedule is *unrecoverable*



27

Aborting a Transaction

28

- If a transaction T_i is aborted, all its actions have to be undone. Not only that, if T_j reads an object last written by T_i , T_j must be aborted as well!
- Most systems try to avoid such *cascading aborts*
 - ▣ If T_i writes an object, T_j can read this only after T_i commits.

28

Aborting a Transaction (cont'd)

29

- In order to *undo* the actions of an aborted transaction, the DBMS maintains a *log* in which every write is recorded.
- This mechanism is also used to recover from system crashes: all active Xacts at the time of the crash are aborted when the system comes back up.

29

Recoverable Schedules and Avoid Cascading Aborts

30

Recoverable

- Aborting T_1 requires aborting T_2
 - But T_2 has already committed!
- A recoverable schedule is one in which this cannot happen.
 - i.e. a Xact commits only after all the Xacts it “depends on” (i.e. it reads from) commit.

No		Yes	
T1	T2	T1	T2
R(A)		R(A)	
W(A)		W(A)	
	R(A)		R(A)
	W(A)		W(A)
Abort	Commit	Commit	Commit

Avoid cascading abort (ACA)

- Aborting T_1 requires aborting T_2 !
- Aborting a Xact can be done without cascading the abort to other Xacts.
- A Xact only reads data from committed Xacts.
- ACA implies recoverable (but not vice-versal!)

T1	T2	T1	T2
R(A)		R(A)	
W(A)		W(A)	
	R(A)		R(A)
	W(A)	Commit	
Abort			R(A)
			W(A)

30

Example

31

W1(X), R2(Y), R1(Y), R2(X), C2, C1

T1, T2: W1(X), R1(Y), R2(Y), R2(X)

- Serializable: Yes, equivalent to T1,T2
- Recoverable: No. Yes, if C1 and C2 are switched
- ACA: No.
 - Yes, if T1 commits before T2 reads X.

31

Including Aborts in Serializability

32

- Extend the definition of a serializable schedule to include aborts
- Serializable schedule: a schedule that is equivalent to some serial execution of the set of *committed* transactions.

32

Conflict Serializable Schedules

33

- Two schedules are *conflict equivalent* if:
 - ▣ Involve the same actions of the same transactions
 - ▣ Every pair of conflicting actions is ordered the same way
- Schedule S is *conflict serializable* if S is conflict equivalent to some serial schedule

33

Recall Conflicts

34

- Two writes by T_i, T_j to same element
 - ▣ $W_i(X); W_j(X)$
- Read/write by T_i, T_j to same element
 - ▣ $W_i(X); R_j(X)$
 - ▣ $R_i(X); W_j(X)$

34

Conflict Equivalent

35

- Outcome of a schedule depends on the order of conflicting operations
- Can interchange non-conflicting ops without changing effect of the schedule
- If two schedules S_1 and S_2 are conflict equivalent then they have the same effect
 - $S_1 \leftrightarrow S_2$ by swapping non-conflicting ops

35

Example Slide

36

36

Precedence Graph Test

37

Is a schedule conflict-serializable ?

Simple test:

- ▣ Build a graph of all transactions T_i
- ▣ Edge from T_i to T_j if T_i comes first, and makes an action that conflicts with one of T_j
- ▣ The test: if the graph has no cycles, then it is conflict serializable !

37

Example 1

38

$R2(A); R1(B); W2(A); R3(A); W1(B); W3(A); R2(B); W2(B)$



This schedule is conflict serializable

38

Example Slide

39

39

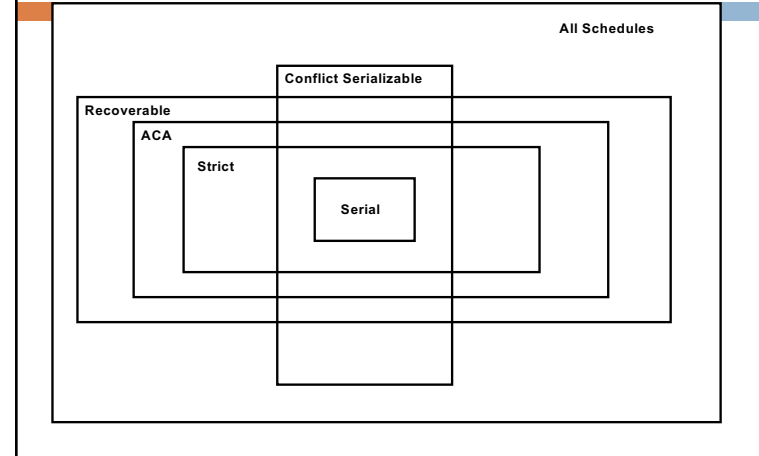
Strict Schedule

40

- A schedule S is strict if a value written by T_i is not read or overwritten by other T_j until T_i aborts or commits
- Example:
 $W1(A); W1(B), C1; W2(A); R2(B); C2;$
- Strict schedules are recoverable, and avoid cascading aborts.

40

Venn Diagram for Schedules



41