

## Performance of Locking

70

- Locking aims to resolve conflicts among transactions by:
    - ▣ Blocking
    - ▣ Aborting
- } Performance penalty
- Blocked Xacts hold locks other Xacts may want
  - Aborting Xact wastes work done thus far
  - Deadlock: Xact is blocked indefinitely until one of the Xacts is aborted

70

## Locking Performance

71

- Locking performance problems are *common!*
- The problem is too much blocking.
- The solution is to reduce the “locking load”
- Good heuristic – If more than 30% of transactions are blocked, then reduce the number of concurrent transactions


Credit: Phil Bernstein

71

## Performance

72

- Locking overhead is primarily due to delays from blocking; minimizes throughput
- What happens to throughput as you increase the # of Xacts?



72

## Improving Performance

73

- Lock the smallest sized object
  - ▣ Reduce likelihood that two Xacts need the same lock
- Reduce the time Xacts hold locks
- Reduce **hot spots**. A hot spot is an object that is frequently accessed, and causes blocking delays

73

## Locking Granularity

74

- **Granularity** - size of data items to lock
  - ▢ e.g., files, pages, records, fields
- Coarse granularity implies
  - ▢ very few locks, so little locking overhead
  - ▢ must lock large chunks of data, so high chance of conflict, so concurrency may be low
- Fine granularity implies
  - ▢ many locks, so high locking overhead
  - ▢ locking conflict occurs only when two transactions try to access the exact same data concurrently

74

## Deadlocks

75


- Deadlock: Cycle of transactions waiting for locks to be released by each other.
- Two ways of dealing with deadlocks:
  - ▢ Deadlock detection
  - ▢ Deadlock prevention

R. Ramakrishnan and J. Gehrke

75

## Deadlocks

76



T1 is waiting for T2 to release its lock  
 T2 is waiting for T1 to release its lock  
 → Such a cycle of transactions is a *deadlock*

### Implications:

- T1 and T2 will make no further progress
- They may hold locks needed by other Xacts
- DBMS tries to prevent or detect (and resolve) deadlocks

76

## Deadlock Detection

77

- Detect deadlocks automatically, and abort a deadlocked transaction (the *victim*).
- Preferred approach, because it allows higher resource utilization
- Timeout-based deadlock detection - If a transaction is blocked for too long, then abort it.
  - ▢ Simple and easy to implement
  - ▢ But aborts unnecessarily (pessimistic) and
  - ▢ some deadlocks persist for too long

77

## Deadlock Detection


78

- Create a **waits-for graph**:
  - ▣ Nodes are transactions
  - ▣ There is an edge from  $T_i$  to  $T_j$  if  $T_i$  is waiting for  $T_j$  to release a lock
  - ▣ Lock mgr adds edge when lock request is queued
  - ▣ Remove edge when lock request granted
- A deadlock exists if there is a cycle in the waits-for graph
- Periodically check for cycles

78

## Waits-For Graph: Example

79



79

## Lecture Example

80

80

## Selecting a Victim

81

- Deadlock is resolved by aborting a Xact in the cycle, and releasing its locks
- Different criteria may be used:

81

## Commercial DBMS Approaches

82

- **MS SQL Server:** Aborts the transaction that is “cheapest” to roll back.
  - ▣ “Cheapest” is determined by the amount of log generated.
  - ▣ Allows transactions that you’ve invested a lot in, to complete.
- **Oracle:** The transaction that detects the deadlock is the victim.
- **DB2:** the deadlock detector arbitrarily selects one deadlocked process as the victim to roll back.

82

## Deadlock Prevention

83

- When there is a high level of lock contention and an increased likelihood of deadlocks
- Prevent deadlocks by giving each Xact a **priority**
- Assign priorities based on timestamps.
  - ▣ Lower timestamp indicates higher priority
  - ▣ i.e., oldest transaction has the highest priority
  - ▣ Higher priority Xacts cannot wait for lower priority Xacts (or vice versa)

83

## Deadlock Prevention (cont’d)

84

- Assume  $T_i$  wants a lock that  $T_j$  holds. Two policies are possible:
  - **Wait-Die:** If  $T_i$  has higher priority,  $T_i$  waits for  $T_j$ ; otherwise  $T_i$  aborts
  - **Wound-wait:** If  $T_i$  has higher priority,  $T_j$  aborts; otherwise  $T_i$  waits
- Both schemes will cause aborts even though deadlock may not have occurred.
- If a transaction re-starts, make sure it has its original timestamp – WHY?

84

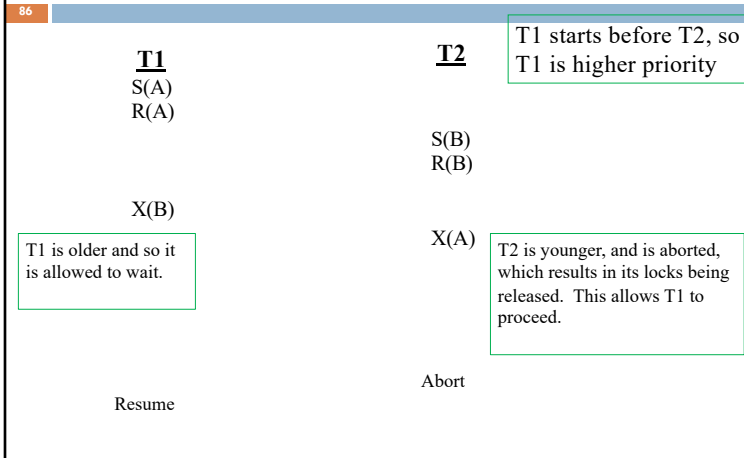
## Wait-Die

85

- Suppose  $T_i$  tries to lock an item already locked by  $T_j$ .
- If  $T_i$  is the older transaction then  $T_i$  will wait
- Otherwise  $T_i$  is aborted and re-starts later with the same timestamp.
- Lower priority transactions never wait for higher priority transactions.

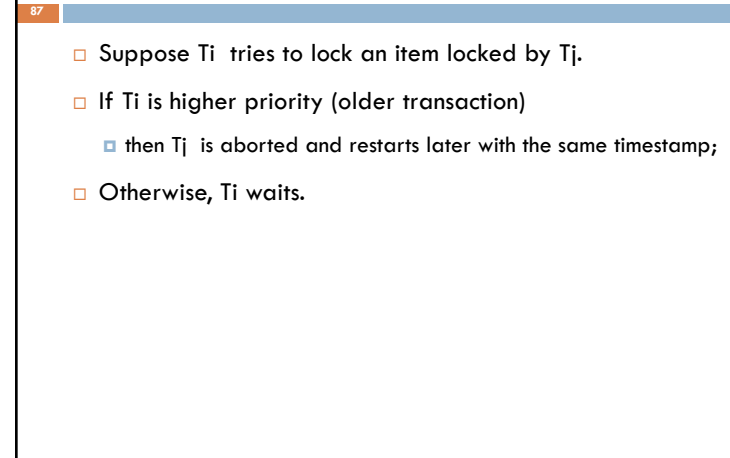
85

## Example: Wait-Die



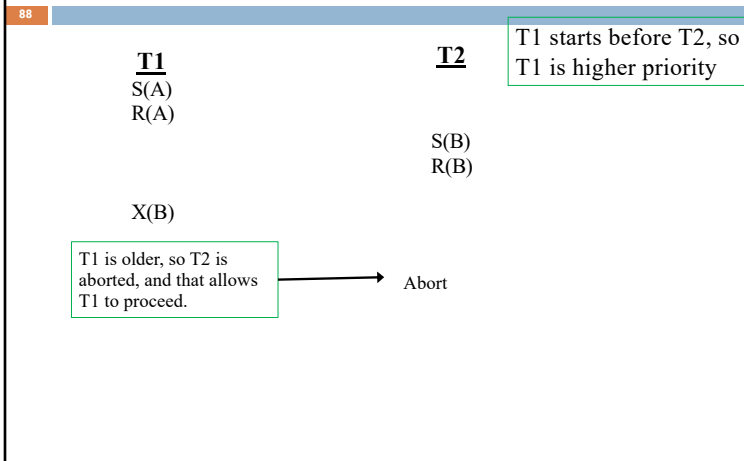
86

## Wound-Wait



87

## Example: Wound-Wait



88