

Announcements

16

- Use **db2srv3** server to access **SE3DB3** database instance

16

Unique Values for Tuples

17

RegNum	Surname	FirstName	BirthDate	DegreeProg
284328	Smith	Luigi	29/04/59	Computing
296328	Smith	John	29/04/59	Computing
587614	Smith	Lucy	01/05/61	Engineering
934856	Black	Lucy	01/05/61	Fine Art
965536	Black	Lucy	05/03/58	Fine Art

- Registration number identifies students, i.e., there is no pair of tuples with the same value for **RegNum**.
- Personal data could identify students as well, i.e., there is no pair of tuples with the same values for all of **Surname**, **FirstName**, **BirthDate**.

17

Keys

18

- A **key** is a set of attributes that uniquely identifies tuples in a relation.
- More precisely:
 - ▢ A set of attributes **K** is a **superkey** for a relation **r** if **r** cannot contain two distinct tuples **t₁** and **t₂** such that **t₁[K]=t₂[K]**;
 - ▢ **K** is a (**candidate**) **key** for **r** if **K** is a minimal superkey (that is, there exists no other superkey **K'** of **r** that is contained in **K** as proper subset, i.e., $K' \subset K$)

18

Example

19

RegNum	Surname	FirstName	BirthDate	DegreeProg
284328	Smith	Luigi	29/04/59	Computing
296328	Smith	John	29/04/59	Computing
587614	Smith	Lucy	01/05/61	Engineering
934856	Black	Lucy	01/05/61	Fine Art
965536	Black	Lucy	05/03/58	Fine Art

- **RegNum** is a key: i.e., **RegNum** is a superkey and it contains a sole attribute, so it is minimal.
- **{Surname, FirstName, BirthDate}** is another key

19

Beware!

20

RegNum	Surname	FirstName	BirthDate	DegreeProg
296328	Smith	John	29/04/59	Computing
587614	Smith	Lucy	01/05/61	Engineering
934856	Black	Lucy	01/05/61	Fine Art
965536	Black	Lucy	05/03/58	Engineering

- There is no pair of tuples with the same values on both **Surname** and **DegreeProg**;

i.e., in each program students have different surnames; can we conclude that **Surname** and **DegreeProg** form a key for this relation?

No! There **could be** students with the same surname in the same program

20

Existence of Keys

21

- Relations are sets; therefore each relation is composed of distinct tuples.
- It follows that the whole set of attributes for a relation defines a **superkey**.
- Therefore **each relation has a key**, which is the set of all its attributes, or a subset thereof.
- The existence of keys guarantees that each piece of data in the database can be accessed,
- Keys are a major feature of the Relational Model and allow us to say that it is "**value-based**".

21

Keys and Null Values

22

If there are nulls, keys do not work that well:

- They do not guarantee unique identification;
- They do not help in establishing correspondences between data in different relations

RegNum	Surname	FirstName	BirthDate	DegreeProg
NULL	Smith	John	NULL	Computing
587614	Smith	Lucy	01/05/61	Engineering
934856	Black	Lucy	NULL	NULL
NULL	Black	Lucy	05/03/58	Engineering

- Are the third and fourth tuple the same?
- How do we access the first tuple?

22

Primary Keys

23

- The presence of nulls in keys has to be limited.
- Each relation must have a **primary key** on which nulls are not allowed (in any attribute)
- Notation: the attributes of the primary key are underlined
- References between relations are realized through primary keys

<u>RegNum</u>	<u>Surname</u>	<u>FirstName</u>	<u>BirthDate</u>	<u>DegreeProg</u>
643976	Smith	John	NULL	Computing
587614	Smith	Lucy	01/05/61	Engineering
934856	Black	Lucy	NULL	NULL
735591	Black	Lucy	05/03/58	Engineering

23

Do we Always Have Primary Keys?

24

- In most cases, we do have reasonable primary keys (e.g., student number, SIN)
- There may be multiple keys, one of which is designated as primary.

24

Recap

25

- A set of fields is a **key** for a relation if:
 1. No two distinct tuples can have same values in all key fields, and
 2. This is not true for any subset of the key.
- If #2 false, then a **superkey**.
- If there's >1 key for a relation, one of the keys is chosen to be the **primary key**.
- E.g., *sid* is a key for Students. (What about *name*?) The set {*sid*, *gpa*} is a superkey.

25

Primary and Candidate Keys

26

- | | |
|--|--|
| <ol style="list-style-type: none"> 1. "For a given student and course, there is a single grade." vs. 2. "Students can take only one course, and receive a single grade for that course; further, no two students in a course receive the same grade." <ul style="list-style-type: none"> ▪ Be careful to define Integrity Constraints (ICs) correctly at design time. ▪ ICS are checked when data is updated. | <p>Enrolled(<i>sid</i>, <i>cid</i>, <i>grade</i>)</p> <p>Enrolled(<u><i>sid</i></u>, <u><i>cid</i></u>, <i>grade</i>)</p> <p>Enrolled(<u><i>sid</i></u>, <i>cid</i>, <i>grade</i>)</p> <ul style="list-style-type: none"> • key (<i>cid</i>, <i>grade</i>) |
|--|--|

26

Foreign Keys

27

- Pieces of data in different relations are correlated by means of values of primary keys.
- Referential integrity constraints are imposed in order to guarantee that the values refer to existing tuples in the referenced relation.
- A **foreign key** requires that the values on a set X of attributes of a relation R₁ must appear as values for the primary key of another relation R₂.
 - In other words, set of attributes in one relation that is used to 'refer' to a tuple in another relation. (Must correspond to primary key of the second relation.) Like a 'logical pointer'.

27

Referential Integrity

28

- E.g. *sid* is a foreign key referring to **Students**:
 - ▢ Enrolled(*sid*: string, cid: string, grade: string)
 - ▢ If all foreign key constraints are enforced, referential integrity is achieved, i.e., no dangling references.

28

Referential Integrity (cont'd)

29

- Only students listed in the Students relation should be allowed to enroll for courses.

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

29

Enforcing Referential Integrity

30

- Consider Students and Enrolled; *sid* in Enrolled is a foreign key that references Students.
- What should be done if an Enrolled tuple with a non-existent student id is inserted? Reject it!
- What should be done if a Students tuple is deleted?
 - ▢ Also delete all Enrolled tuples that refer to it.
 - ▢ Disallow deletion of a Students tuple that is referred to.
 - ▢ Set *sid* in Enrolled tuples that refer to it to a *default sid*.
 - ▢ Set *sid* in Enrolled tuples that refer to it to NULL.
- Similar if primary key of Students tuple is updated.

30

Where do ICs Come From?

31

- ICs are based upon the semantics of the real-world enterprise that is being described in the database relations.
- We can check a database instance to see if an IC is violated, but we cannot infer that an IC is true by looking at an instance.
 - ▢ An IC is a statement about *all possible* instances
- Key and foreign key ICs are the most common; more general ICs supported too.

31

One More Example

32

Offences	<u>Code</u>	Date	Officer	Dept	Registration
	143256	25/10/1992	567	75	5694 FR
	987554	26/10/1992	456	75	5694 FR
	987557	26/10/1992	456	75	6544 XY
	630876	15/10/1992	456	47	6544 XY
	539856	12/10/1992	567	47	6544 XY

Officers	<u>RegNum</u>	Surname	FirstName
	567	Brun	Jean
	456	Larue	Henri
	638	Larue	Jacques

Cars	<u>Registration</u>	<u>Dept</u>	Owner
	6544 XY	75	Cordon Edouard
	7122 HT	75	Cordon Edouard
	5694 FR	75	Latour Hortense
	6544 XY	47	Mimault Bernard

- $\text{Offences}[\text{Officer}] \subseteq \text{Officers}[\text{RegNum}]$
- $\text{Offences}[\text{Registration}, \text{Dept}] \subseteq \text{Cars}[\text{Registration}, \text{Dept}]$

32

Violation of Foreign keys

33

Offences	<u>Code</u>	Date	Officer	Dept	Registration
	987554	26/10/1992	456	75	5694 FR
	630876	15/10/1992	456	47	6544 XY

Officers	<u>RegNum</u>	Surname	FirstName
	567	Brun	Jean
	638	Larue	Jacques

Cars	<u>Registration</u>	<u>Dept</u>	Owner	...
	7122 HT	75	Cordon Edouard	...
	5694 FR	93	Latour Hortense	...
	6544 XY	47	Mimault Bernard	...

33