

## Example: =ANY

1		
Sells		
Bar	Beer	Price
Jane	Miller	3.00
Joe	Miller	4.00
Joe	Bud	3.00
Jack	Bud	4.00
Tom	Miller	4.50

Result		
Bar		
Jane		
Joe		

```
SELECT Bar
FROM Sells
WHERE Beer = 'Miller' AND Price =
  ANY(SELECT Price
      FROM Sells
      WHERE Beer='Bud')
```

1

## The Exists Operator

- EXISTS(<subquery>) is true if and only if the subquery result is not empty.
- Example: From Beers(name, manf), find those beers that are the unique (only) beer made by their manufacturer.

Credit: Renee J. Miller

2

## Example: EXISTS

3		
<p>SELECT name FROM Beers b1 WHERE NOT EXISTS (</p> <div> <p>SELECT *</p> <p>FROM Beers</p> <p>WHERE manf = b1.manf AND name &lt;&gt; b1.name);</p> </div> <p>Set of beers with the same manf as b1, but not the same beer</p> <p>Notice scope rule: manf refers to closest nested FROM with a relation having that attribute. (Some DBMS consider this ambiguous.)</p> <p>Notice the SQL "not equals" operator</p>		


3

## Union, Intersection, and Difference

- Union, intersection, and difference of relations are expressed by the following forms, each involving subqueries:
  - (<subquery>) UNION (<subquery>)
  - (<subquery>) INTERSECT (<subquery>)
  - (<subquery>) EXCEPT (<subquery>)

4

## Visually



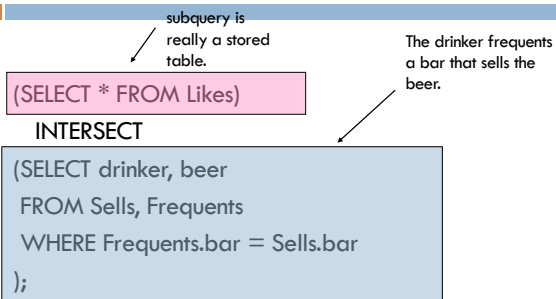
5

## Example: Intersection

- Using `Likes(drinker, beer)`, `Sells(bar, beer, price)`, and `Frequents(drinker, bar)`, find the drinkers and beers such that:
  - The drinker likes the beer, and
  - The drinker frequents at least one bar that sells the beer.

6

## Solution



7

## Bag Semantics

- A *bag* (or *multiset*) is like a set, but an element may appear more than once.
- Example:  $\{1,2,1,3\}$  is a bag.
- Example:  $\{1,2,3\}$  is also a bag that happens to be a set.

8

## Bag (Multiset) Semantics

9

- SQL primarily uses bag semantics
- The SELECT-FROM-WHERE statement uses bag semantics
  - ▢ originally for efficiency reasons
- The default for union, intersection, and difference is set semantics.
  - ▢ That is, duplicates are eliminated as the operation is applied.

9

## Motivation: Efficiency

10

- When doing projection, it is easier to avoid eliminating duplicates.
  - ▢ Just work tuple-at-a-time.
- For intersection or difference, it is most efficient to sort the relations first.
  - ▢ At that point you may as well eliminate the duplicates anyway.

10

## Controlling Duplicate Elimination

11

- Force the result to be a set by SELECT DISTINCT . . .
- Force the result to be a bag (i.e., don't eliminate duplicates) by ALL, as in . . . UNION ALL . . .

11

## Example: DISTINCT

12

- From `Sells(bar, beer, price)`, find all the different prices charged for beers:

```
SELECT DISTINCT price
FROM Sells;
```

Notice that without DISTINCT, each price would be listed as many times as there were bar/beer pairs at that price.

12

## Example: ALL

13

- Using relations **Frequents(drinker, bar)** and **Likes(drinker, beer)**:
- Lists drinkers who frequent more bars than they like beers, and do so as many times as the difference of those counts.

```
(SELECT drinker FROM Frequents)
  EXCEPT ALL
(SELECT drinker FROM Likes);
```

13

## Ordering the Display of Tuples

14

- List in alphabetic order the names of all instructors
 

```
select name
from   instructor
order by name
```
- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.
  - Example: **order by name desc**

Credit: Silberchatz, Korth &amp; Sudarshan

14

## Humour

15

SQL query walks into a bar, and approaches two tables and asks, can I join you?



15