

Cross Product

1

- Evaluating joins involves combining two or more relations
- Given two relations, S and R, each row of S is paired with each row of R
- Result schema: one attribute from each attribute of S and R

Example

2

Sells

Bar	Beer	Price
Joe	Bud	3.00
Tom	Miller	4.00
Jane	Lite	3.25

Frequents

Drinker	Bar
Aaron	Joe
Mary	Jane

Cross product, also known as the Cartesian product

Sells x Frequents

(Bar)	Beer	Price	Drinker	(Bar)
Joe	Bud	3.00	Aaron	Joe
Joe	Bud	3.00	Mary	Jane
Tom	Miller	4.00	Aaron	Joe
Tom	Miller	4.00	Mary	Jane
Jane	Lite	3.25	Aaron	Joe
Jane	Lite	3.25	Mary	Jane

SELECT drinker
FROM Frequents, Sells
WHERE beer = 'Bud' AND
Frequents.bar = Sells.bar

Drinker
Aaron

2

Joined Relations

3

- **Join operations** take two relations and return as a result another relation.
- A join operation is a Cartesian product which requires that tuples in the two relations match (under some condition). It also specifies the attributes that are present in the result of the join

©Silberschatz, Korth and Sudarshan

3

Join Operations – Example

4

- Relation course

course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

- Relation prereq

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

- Observe that

prereq information is missing for CS-315 and

course information is missing for CS-347

4

Outer Join

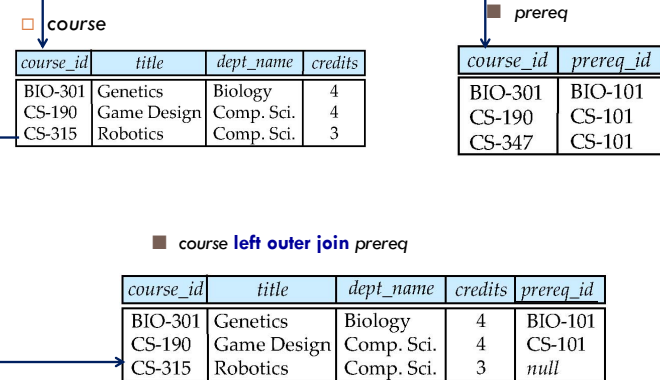
5

- An extension of the join operation that avoids loss of information.
- Suppose you have two relations R and S. A tuple of R that has no tuple of S with which it joins is said to be **dangling**.
 - ▣ Similarly for a tuple of S.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Outerjoin preserves dangling tuples by padding them with NULL.

5

Left Outer Join

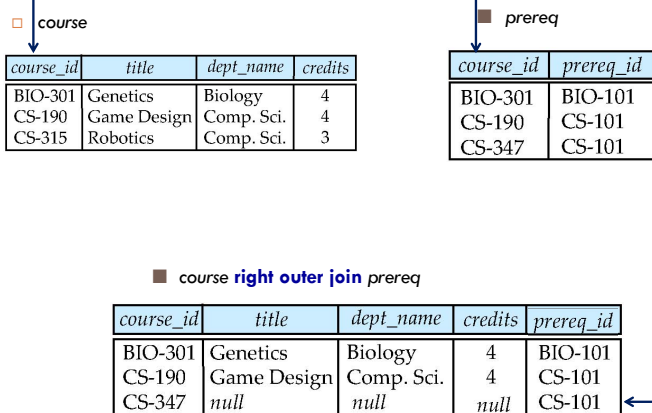
6



6

Right Outer Join

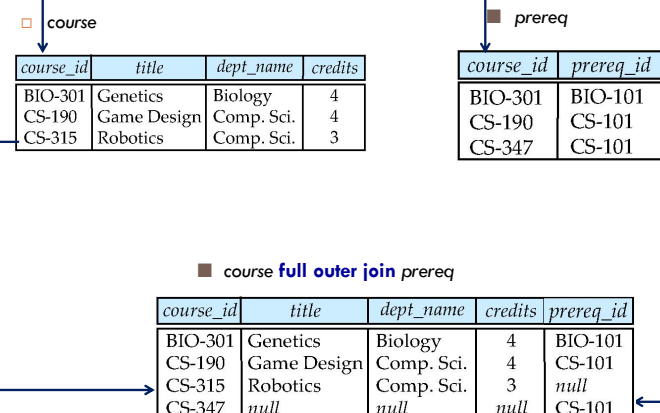
7



7

Full Outer Join

8



8

Inner Join

9

□ course ■ prereq

course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

□ course **inner join** prereq **on**
course.course_id = prereq.course_id

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101

9

Outerjoins

- 10
- R **OUTER JOIN** S is the core of an outerjoin expression. It is modified by:
1. Optional **NATURAL** in front of **OUTER**.
 - Check equality on all common attributes
 - No two attributes with the same name in the output
 2. Optional **ON <condition>** after **JOIN**.
 3. Optional **LEFT**, **RIGHT**, or **FULL** before **OUTER**.
 - ◆ **LEFT** = pad dangling tuples of R only.
 - ◆ **RIGHT** = pad dangling tuples of S only.
 - ◆ **FULL** = pad both; this choice is the default.

Credit: Renee J. Miller

10

Class Example

11

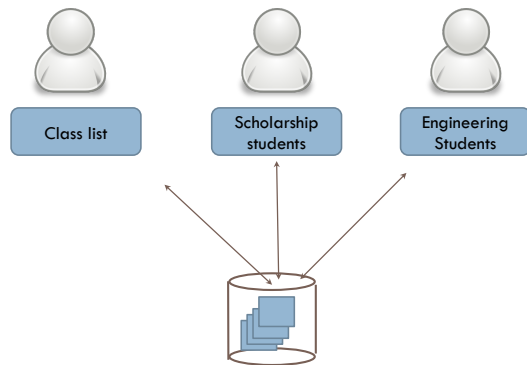
11

VIEWS

12

Scenario

13



13

Views

14

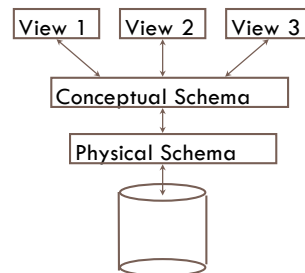
- In most cases, it is not desirable for all users to see the entire data instance.
- A **view** provides a mechanism to hide certain data from the view of certain users.

14

Levels of Abstraction

15

- Many views, single conceptual (logical) schema and physical schema.
 - ▣ Views describe how users see the data.
 - ▣ Conceptual schema defines logical structure
 - ▣ Physical schema describes the files and indexes used.



Credit: Renee J. Miller

15

Views

16

- A **view** is a relation defined in terms of stored tables (called **base tables**) and other views.
- Two kinds:
 1. **Virtual** = not stored in the database; just a query for constructing the relation.
 2. **Materialized** = actually constructed and stored.

16

Declaring Views

17

- Declare by:

```
CREATE [MATERIALIZED] VIEW <name> AS <query>;
```

- A view name
- A possible list of attribute names (for example, when arithmetic operations are specified or when we want the names to be different from the attributes in the base relations)
- A query to specify the view contents
- Default is virtual.

17

Example: View Definition

18

- **CanDrink(drinker, beer)** is a view “containing” the drinker-beer pairs such that the drinker frequents at least one bar that serves the beer:

```
CREATE VIEW CanDrink AS
  SELECT drinker, beer
  FROM Frequents, Sells
  WHERE Frequents.bar = Sells.bar;
```

18

Example: Accessing a View

19

- Query a view as if it were a base table.
 - ▣ Also: a limited ability to modify views if it makes sense as a modification of one underlying base table.
- **Example query:**

```
SELECT beer FROM CanDrink
WHERE drinker = 'Sally';
```

19

Another Example

20

- **Example:** View Synergy has **(drinker, beer, bar)** triples such that the bar serves the beer, the drinker frequents the bar and likes the beer.

20

Example: The View

21

CREATE VIEW Synergy AS

SELECT Likes.drinker, Likes.beer, Sells.bar

FROM Likes, Sells, Frequents

WHERE Likes.drinker = Frequents.drinker

AND Likes.beer = Sells.beer

AND Sells.bar = Frequents.bar;

Natural join of Likes,
Sells, and Frequents

Pick one copy of
each attribute

21

Updates on Views

22

- Generally, it is impossible to modify a virtual view, because it doesn't exist.
- Can't we "translate" updates on views into "equivalent" updates on base tables?
 - ▣ Not always (in fact, not often)
 - ▣ Most systems prohibit most view updates
- We cannot insert into Synergy --- it is a virtual view.

22

Interpreting a View Insertion

23

- But we could try to translate a (drinker, beer, bar) triple into three insertions of projected pairs, one for each of Likes, Sells, and Frequents.

23

Insertion

24

```
INSERT INTO LIKES VALUES(n.drinker, n.beer);
INSERT INTO SELLS(bar, beer) VALUES(n.bar, n.beer);
INSERT INTO FREQUENTS VALUES(n.drinker, n.bar);
```

- ▣ Sells.price will have to be NULL.
- ▣ There isn't always a unique translation.

24

Materialized Views

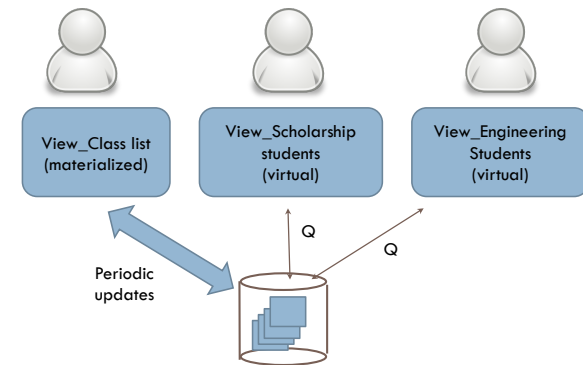
25

- **Materialized** = actually constructed and stored (keeping a temporary table)
- **Concerns:** maintaining correspondence between the base table and the view when the base table is updated
- **Strategy:** incremental update

25

Example

26



26

Example: Class Mailing List

27

- The class mailing list `db3students` is in effect a materialized view of the class enrollment
- Updated periodically
 - ▣ You can enroll and miss an email sent out after you enroll.
- Insertion into materialized view normally followed by insertion into base table

27

Materialized View Updates

28

- Update on a single view without aggregate operations: update may map to an update on the underlying base table (most SQL implementations)
- Views involving joins: an update *may map to an* update on the underlying base relations not always possible

28

Example: A Data Warehouse

29

- Wal-Mart stores every sale at every store in a database.
- Overnight, the sales for the day are used to update a *data warehouse* = materialized views of the sales.
- The warehouse is used by analysts to predict trends and move goods to where they are selling best.

29