

Materialized Views

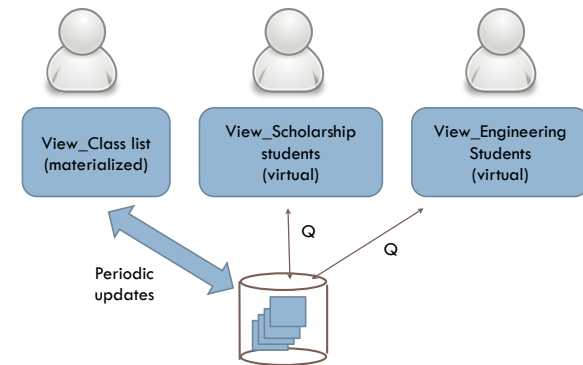
1

- **Materialized** = actually constructed and stored (keeping a temporary table)
- **Concerns:** maintaining correspondence between the base table and the view when the base table is updated
- **Strategy:** incremental update

1

Example

2



2

Example: Class Mailing List

3

- The class mailing list `db3students` is in effect a materialized view of the class enrollment
- Updated periodically
 - ▣ You can enroll and miss an email sent out after you enroll.
- Insertion into materialized view normally followed by insertion into base table

3

Materialized View Updates

4

- Update on a single view without aggregate operations: update may map to an update on the underlying base table (most SQL implementations)
- Views involving joins: an update *may map to an* update on the underlying base relations not always possible

4

Example: A Data Warehouse

5

- Wal-Mart stores every sale at every store in a database.
- Overnight, the sales for the day are used to update a *data warehouse* = materialized views of the sales.
- The warehouse is used by analysts to predict trends and move goods to where they are selling best.

5

INDEXES

6

Example

7

- Find the price of beers manufactured by Pete's and sold by Joe.

```
SELECT price
FROM Beers, Sells
WHERE manf = 'Pete's' AND bar = 'Joe' AND
Sells.beer = Beers.name
```

7

An Index

8


- A data structure used to speed access to tuples of a relation, based on values of one or more attributes ("search key" fields)
- Organizes records via trees or hashing
- Given a value v , the index takes us to only those tuples that have v in the attribute(s) of the index.
- **Example:** use *BeerInd* (on manf) and *SellInd* (on bar, beer) to find the prices of beers manufactured by Pete's and sold by Joe.

8

B+ Tree Index

9

- The B+ tree structure is the most common index type in databases
- Index files can be quite large, often stored on disk, partially loaded into memory as needed
- Each node is at least 50% full

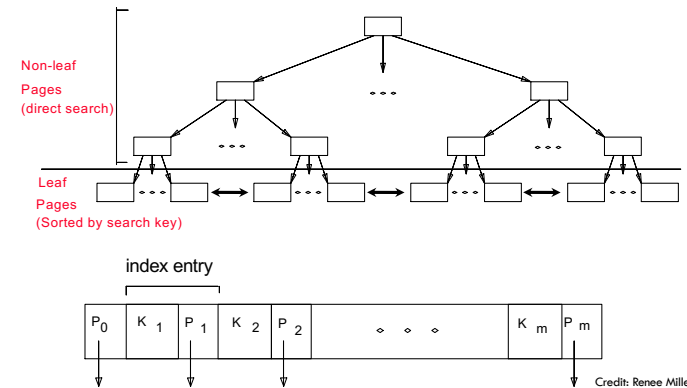


9

B+ Tree Index

10


Supports equality and range-searches efficiently



10

Example

11



- Find 28*? 29*? All > 15* and < 30*
- Insert/delete: Find data entry in leaf, then change it. Need to adjust parent sometimes.
 - And change sometimes bubbles up the tree


11

Inserting a Data Entry

12

- Find correct leaf L.
- Put data entry onto L.
 - If L has enough space, done!
 - Else, must **split** L (into L and a new node L2)
 - Redistribute entries evenly, **copy up** middle key.
 - Insert index entry pointing to L2 into parent of L.
- This can happen recursively
 - To split index node, redistribute entries evenly, but **push up** middle key.
- Splits "grow" tree; root split increases height.

Insert data
value 4




12

Deleting a Data Entry

Delete value 3

13

- ❑ Start at root, find leaf L where entry belongs.
- ❑ Remove the entry.
 - ❑ If L is at least half-full, done!
 - ❑ If not,
 - ❑ Try to **re-distribute**, borrowing from **sibling** (adjacent node with same parent as L).
 - ❑ If re-distribution fails, **merge** L and sibling.
- ❑ If merge occurred, must delete entry (pointing to L or sibling) from parent of L.
- ❑ Merge could propagate to root, decreasing height.




13

Balanced vs. Unbalanced Trees


14

- ❑ In a balanced tree, every path from the root to a leaf node is the same length.

o Balanced



o Unbalanced



Credit: S. Lee

14

Hash Based Indexes

15

- ❑ Good for equality selections.
- ❑ Index is a collection of **buckets**
 - ❑ Bucket = **primary page** plus zero or more **overflow pages**.
 - ❑ Buckets contain data entries.
- ❑ **Hashing function h**: $h(r)$ = bucket in which (data entry for) record r belongs. h looks at the **search key** fields of r.
 - ❑ No need for "index entries" in this scheme.

15

Index Classification


16

- ❑ **Primary vs. secondary**: If search key contains primary key, then called primary index.
 - ❑ **Unique** index: Search key contains a candidate key.
- ❑ **Clustered vs. unclustered**: If order of index data entries is the same as order of data records, then called clustered index.
 - A table can have at most one clustered index – why?

16

Clustered vs. Unclustered Index

17



17

Declaring Indexes

18

□ No standard!

□ Typical syntax:

```
CREATE INDEX BeerInd ON Beers (manf);
CREATE INDEX SellInd ON Sells (bar,
    beer);
```

18

Using Indexes

19

- Given a value v , the index takes us to only those tuples that have v in the attribute(s) of the index.
- **Example:** use BeerInd and SellInd to find the prices of beers manufactured by Pete's and sold by Joe.

19

Using Indexes --- (2)

20

```
SELECT price
FROM Beers, Sells
WHERE manf = 'Pete''s' AND
      Beers.name = Sells.beer AND
      bar = 'Joe''s Bar';
```

1. Use BeerInd to get all the beers made by Pete's.
2. Then use SellInd to get prices of those beers, with bar = 'Joe's Bar'

20

Understanding the Workload

21

- For each query in the workload:
 - Which relations does it access?
 - Which attributes are retrieved?
 - Which attributes are involved in selection/join conditions? How selective are these conditions likely to be?
- For each update in the workload:
 - The type of update (INSERT/DELETE/UPDATE), and the attributes that are affected.

21

Choice of Indexes

22

- What indexes should we create?
 - Which relations should have indexes? What field(s) should be the search key? Should we build several indexes?
- For each index, what kind of an index should it be?
 - Clustered? Hash/tree?

22

Choice of Indexes (cont'd)

23

- **One approach:** Consider the most important queries in turn. Consider the best plan using the current indexes, and see if a better plan is possible with an additional index.
 - Implies an understanding of how a DBMS evaluates queries and creates **query evaluation plans**.
- Before creating an index, must also consider the impact on updates in the workload!
 - **Trade-off:** Indexes can make queries go faster, updates slower. Require disk space, too.

23

Guidelines

24

- Attributes in WHERE clause are candidates for index keys.
 - Exact match condition suggests hash index.
 - Range query suggests tree index.
- Multi-attribute search keys should be considered when a WHERE clause contains several conditions.
- Try to choose indexes that benefit as many queries as possible.
 - Since only one index can be clustered per relation, choose it based on important queries that would benefit the most from clustering.

24

Examples

25

- B+ tree index on E.age can be used to get qualifying tuples.

```
SELECT E.dno
FROM Emp E
WHERE E.age>40
```

- Equality queries and duplicates:
 - Indexing on E.hobby

```
SELECT E.dno
FROM Emp E
WHERE E.hobby='Stamps'
```

25

Composite Search Keys


26

- **Composite Search Keys:** Search on a combination of fields.

- **Equality query:** Every field value is equal to a constant value. E.g. wrt $\langle \text{sal}, \text{age} \rangle$ index:
 - $\text{age}=20$ and $\text{sal}=75$
- **Range query:**
 - $\text{age}=20$ and $\text{sal} > 10$

- Data entries in index sorted by search key to support range queries.

Examples of composite key indexes



26