

## Index Classification


12

- ❑ **Primary vs. secondary:** If search key contains primary key, then called primary index.
  - ❑ **Unique** index: Search key contains a candidate key.
- ❑ **Clustered vs. unclustered:** If order of index data entries is the same as order of data records, then called clustered index.
  - A table can have at most one clustered index – why?

12

## Clustered vs. Unclustered Index

13



13

## Declaring Indexes

14

- ❑ No standard!
- ❑ **Typical syntax:**

```
CREATE INDEX BeerInd ON Beers(manf);
CREATE INDEX SellInd ON Sells(bar,
    beer);
```

14

## Using Indexes

15

- ❑ Given a value  $v$ , the index takes us to only those tuples that have  $v$  in the attribute(s) of the index.
- ❑ **Example:** use BeerInd and SellInd to find the prices of beers manufactured by Pete's and sold by Joe.

15

## Using Indexes --- (2)

16

```
SELECT price
FROM Beers, Sells
WHERE manf = 'Pete''s' AND
      Beers.name = Sells.beer AND
      bar = 'Joe''s Bar';
```

1. Use BeerInd to get all the beers made by Pete's.
2. Then use SellInd to get prices of those beers, with bar = 'Joe's Bar'

16

## Understanding the Workload

17

- For each query in the workload:
  - Which relations does it access?
  - Which attributes are retrieved?
  - Which attributes are involved in selection/join conditions? How selective are these conditions likely to be?
- For each update in the workload:
  - The type of update (INSERT/DELETE/UPDATE), and the attributes that are affected.

17

## Choice of Indexes

18

- What indexes should we create?
  - Which relations should have indexes? What field(s) should be the search key? Should we build several indexes?
- For each index, what kind of an index should it be?
  - Clustered? Hash/tree?

18

## Choice of Indexes (cont'd)

19

- **One approach:** Consider the most important queries in turn. Consider the best plan using the current indexes, and see if a better plan is possible with an additional index.
  - Implies an understanding of how a DBMS evaluates queries and creates **query evaluation plans**.
- Before creating an index, must also consider the impact on updates in the workload!
  - **Trade-off:** Indexes can make queries go faster, updates slower. Require disk space, too.

19

## Guidelines

20

- ❑ Attributes in WHERE clause are candidates for index keys.
  - Exact match condition suggests hash index.
  - Range query suggests tree index.
- ❑ Multi-attribute search keys should be considered when a WHERE clause contains several conditions.
- ❑ Try to choose indexes that benefit as many queries as possible.
  - ❑ Since only one index can be clustered per relation, choose it based on important queries that would benefit the most from clustering.

20

## Examples

21

- ❑ B+ tree index on E.age can be used to get qualifying tuples.

```
SELECT E.dno
FROM Emp E
WHERE E.age>40
```

- ❑ Equality queries and duplicates:

- ❑ Indexing on E.hobby

```
SELECT E.dno
FROM Emp E
WHERE E.hobby='Stamps'
```


21

## Composite Search Keys

22

- ❑ **Composite Search Keys:** Search on a combination of fields.
  - **Equality query:** Every field value is equal to a constant value. E.g. wrt <sal,age> index:
    - age=20 and sal =75
  - **Range query:**
    - age=20 and sal > 10
- ❑ Data entries in index sorted by search key to support range queries.

Examples of composite key indexes



22

## Database Tuning

23

- ❑ A major problem in making a database run fast is deciding which indexes to create.
- ❑ **Pro:** An index speeds up queries that can use it.
- ❑ **Con:** An index slows down all modifications on its relation because the index must be modified too.

23

## Example: Tuning

24

- Suppose the only things we did with our beers database was:
  1. Insert new beers into a relation (10%).
  2. Find the price of a given beer at a given bar (90%).
- Then **SellInd** on Sells(bar, beer) would be helpful, but **BeerInd** on Beers(manf) would be harmful.

24

## Tuning Advisors

25


- A major research area
  - ▣ Because hand tuning is so hard.
- An advisor gets a **query load**, e.g.:
  1. Choose random queries from the history of queries run on the database, or
  2. Designer provides a sample workload.

25

## Tuning Advisors --- (2)

26

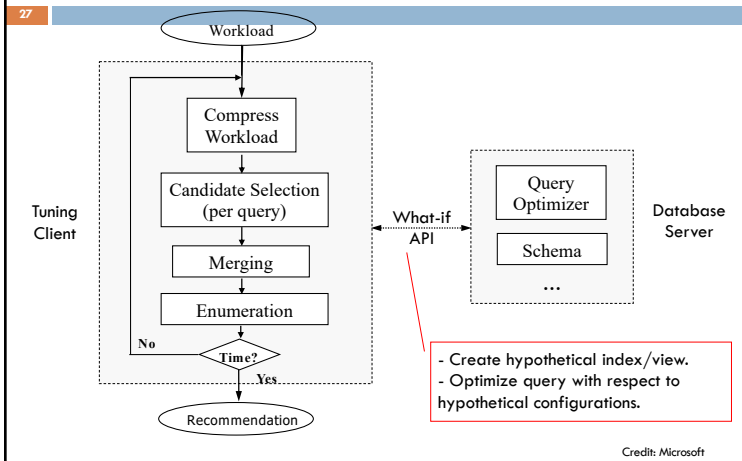
- The advisor generates candidate indexes and evaluates each on the workload.
  - ▣ Measure the improvement/degradation in the average running time of the queries.



26

## Example: Database Tuning Architecture

27



27

## Summary

28

- If selection queries are frequent, sorting the file or building an index is important.
  - Hash-based indexes only good for equality search.
  - Tree-based indexes best for range search; also good for equality search.

28

## Summary (cont'd)

29

- Can have several indexes on a given file of data records, each with a different search key.
- Understanding the nature of the workload for the application
  - What are the important queries and updates? What attributes/relations are involved?
- Indexes are chosen to speed up important queries (and perhaps some updates!).
  - Consider index maintenance overhead on updates to key fields.
  - Choose indexes that can help many queries, if possible.

29